

---

# Automated diagnosis without predictability is a recipe for failure

Raja Sambasivan

Greg Ganger

PARALLEL DATA LABORATORY  
Carnegie Mellon University

# What is important for design? (I)

---

*The designer usually finds himself floundering in a sea of possibilities, unclear about how one choice will limit his freedom...or affect the size and performance of the entire system.*

--*Butler Lampson*

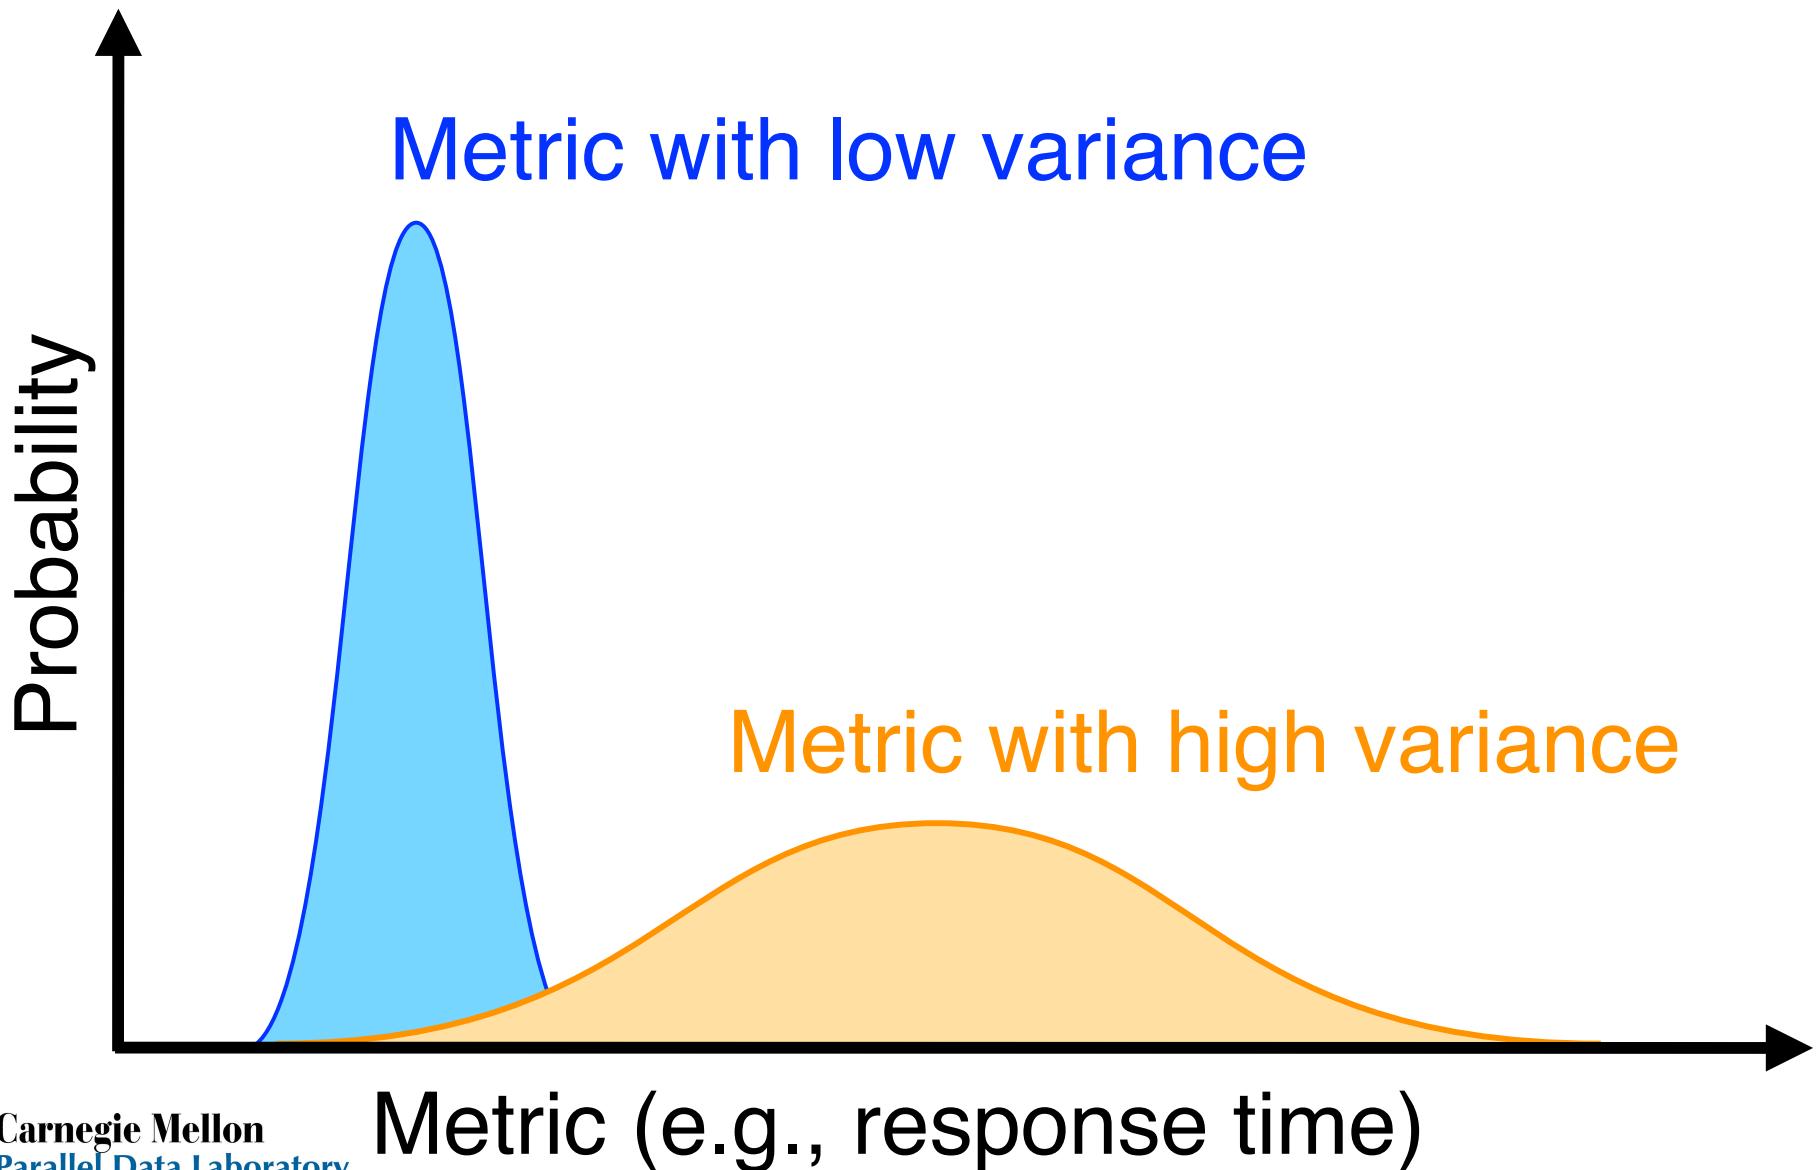
# What is important for design? (II)

---

- Commonly agreed upon axes:
  - Correctness, perf., reliability, power...
- This talk: Predictability also important
  - Especially for large distributed systems
  - Affects ability to optimize other axes
  - Affects manageability

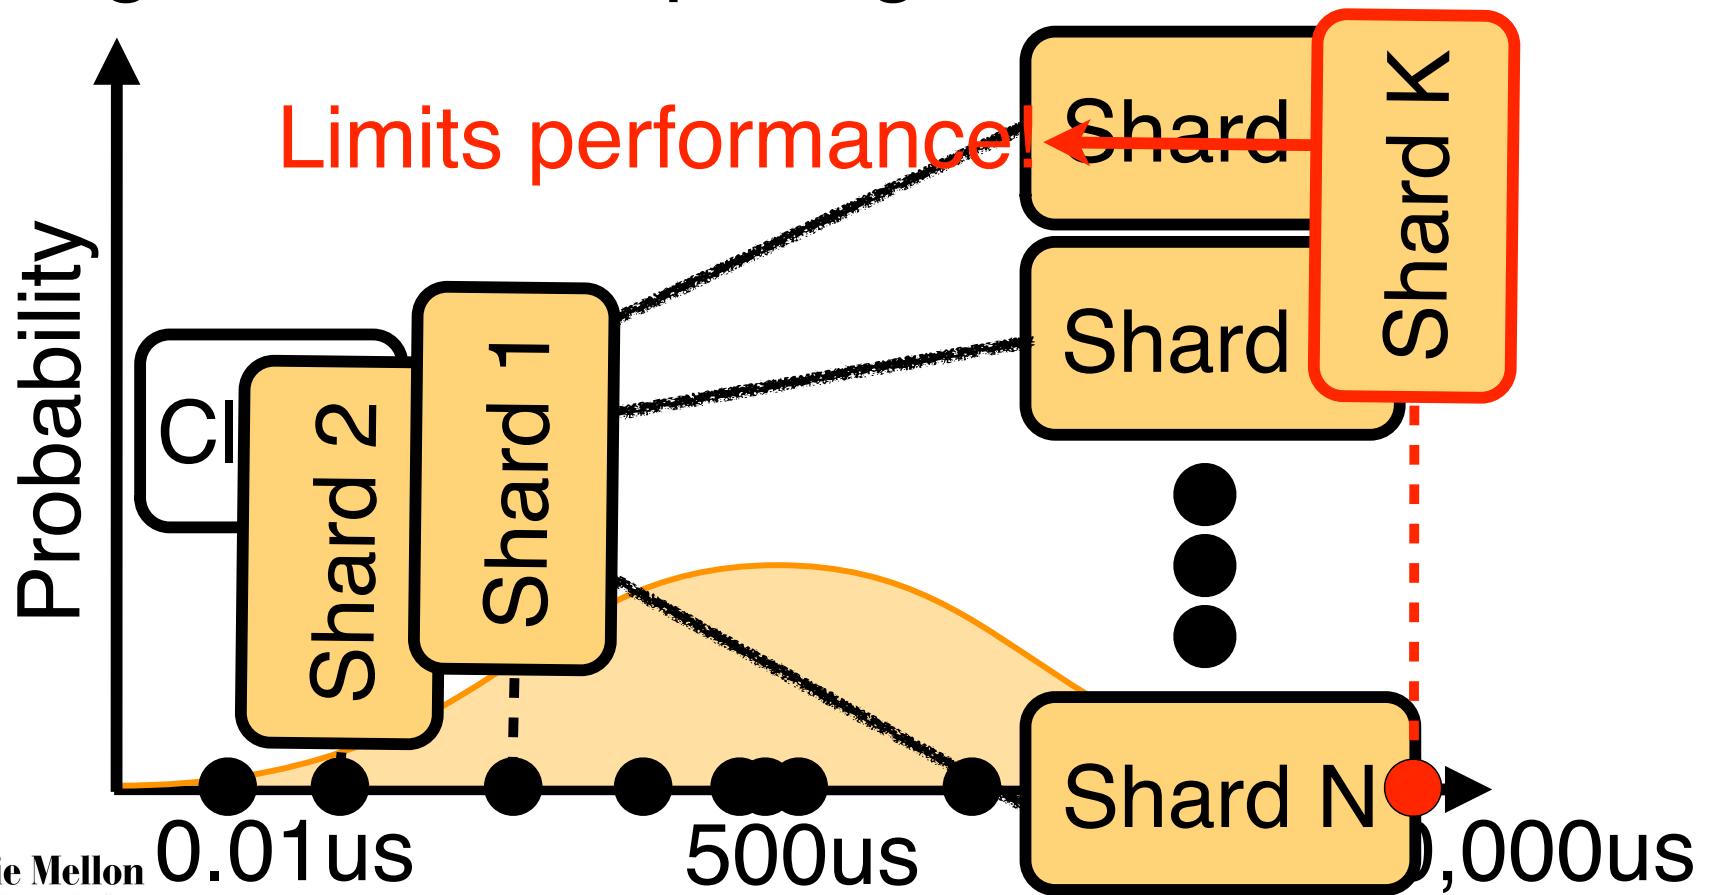
Achieving it will require hard work

# Predictability means low variance



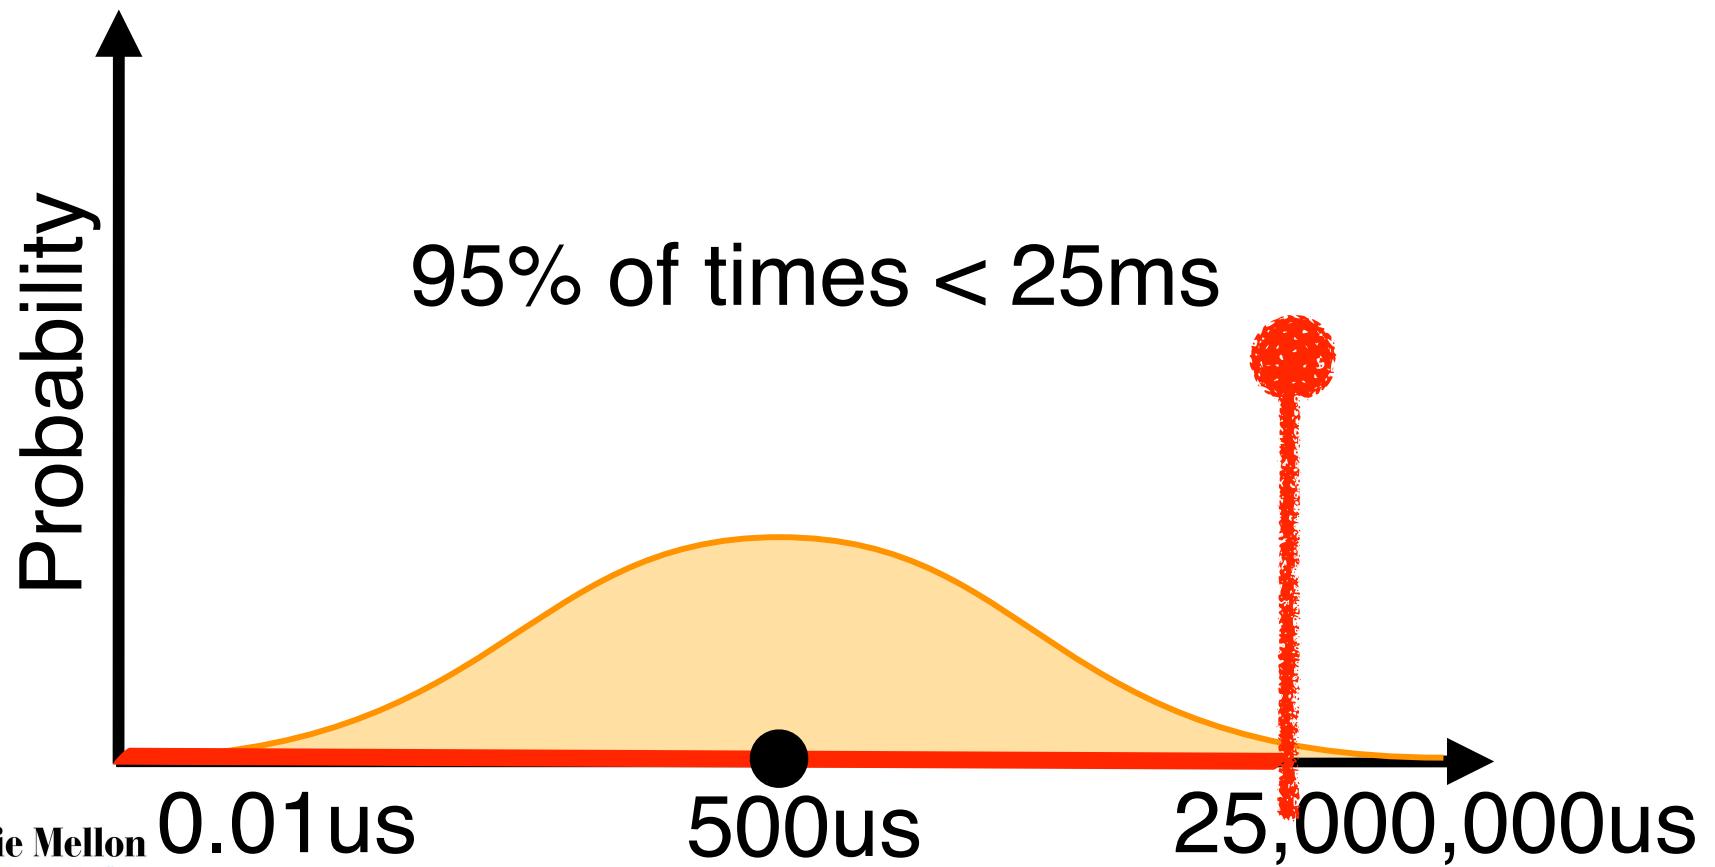
# Predictability is important (I)

- For performance and efficiency of very large-scale computing workloads



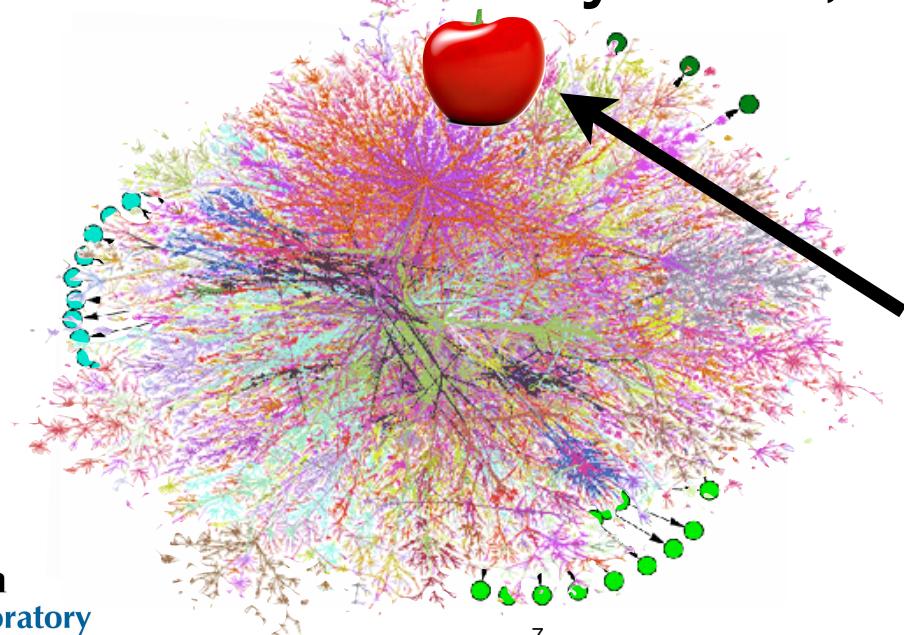
# Predictability is important (II)

- For setting & enforcing SLAs and QoS



# Predictability is important (III)

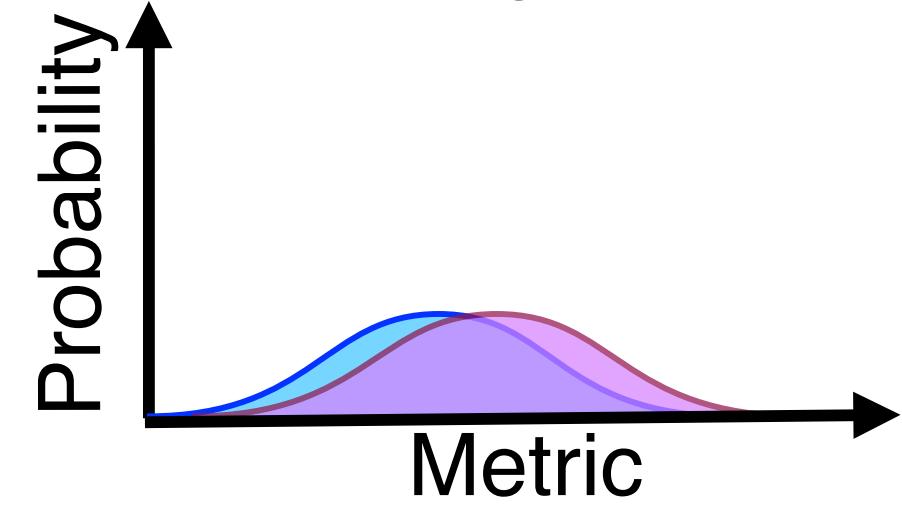
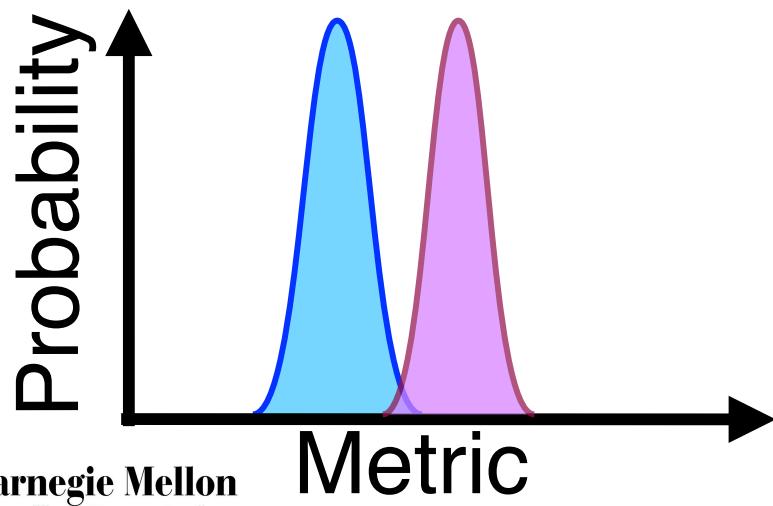
- For success of automated diagnosis  
Magpie | NetMedic | Pinpoint | Spectroscope | ...
- Automation tools not used in production
  - Problem is the system, not the tools...



Cherry-on-top  
approach to  
automation

# The issue with diagnosis tools

- Use deviations in metrics to localize perf. problems to a few likely culprits
- Increased variance in metrics yields...
  - more false positives or false negatives



# Grad student vs. diagnosis tool

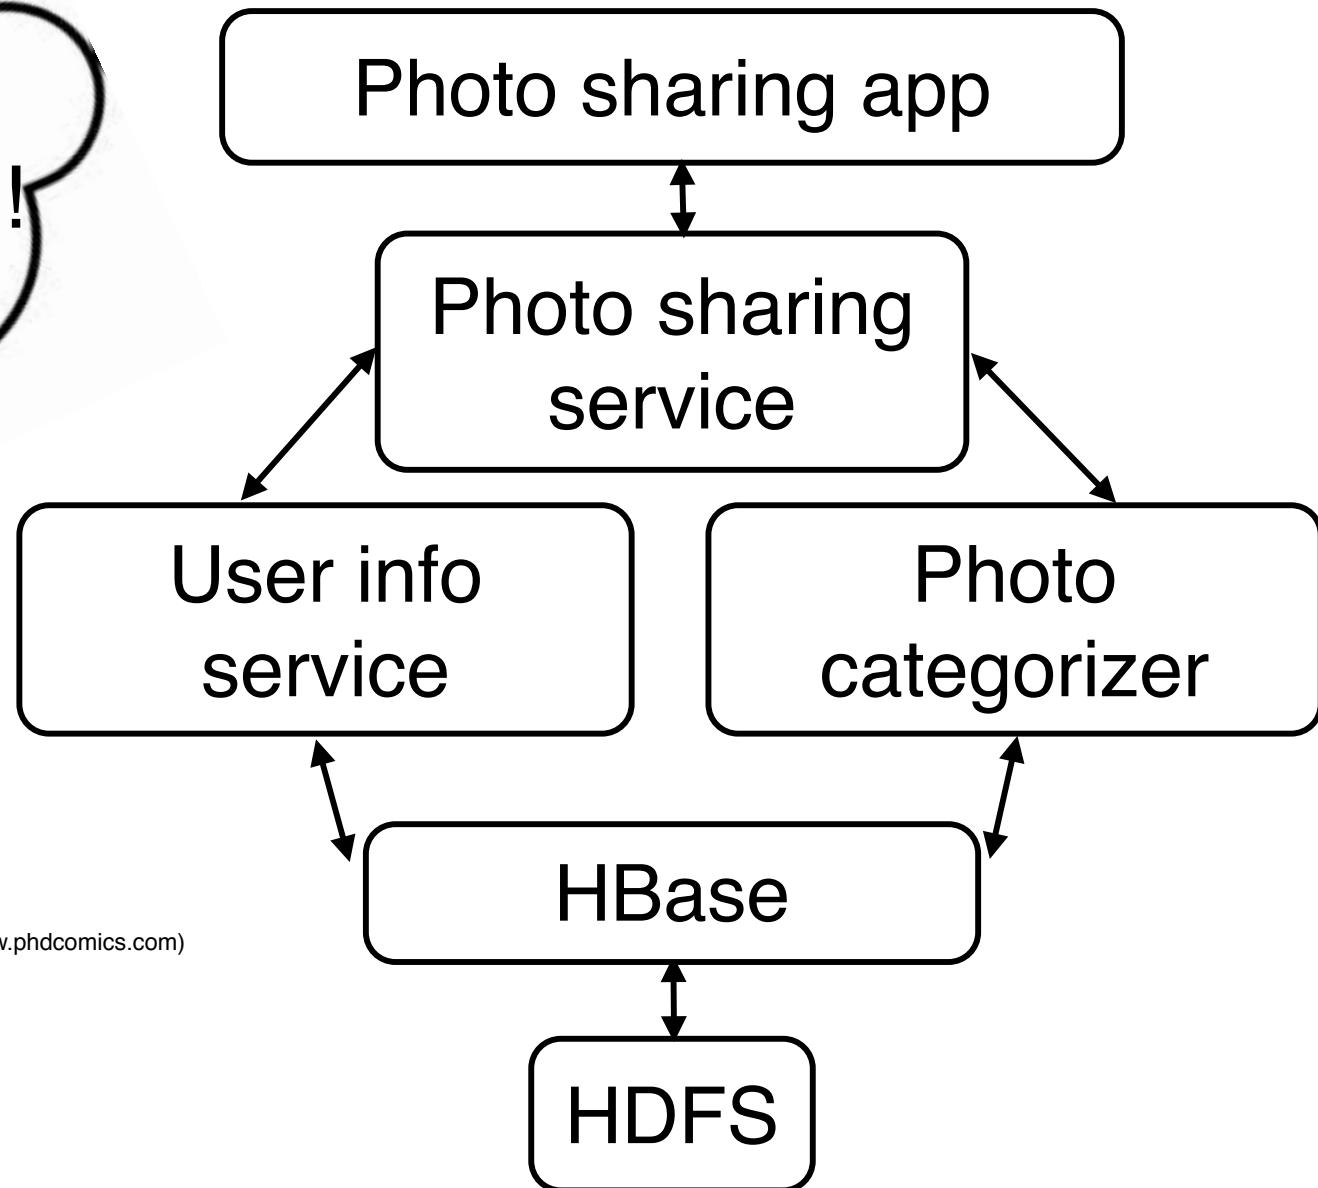
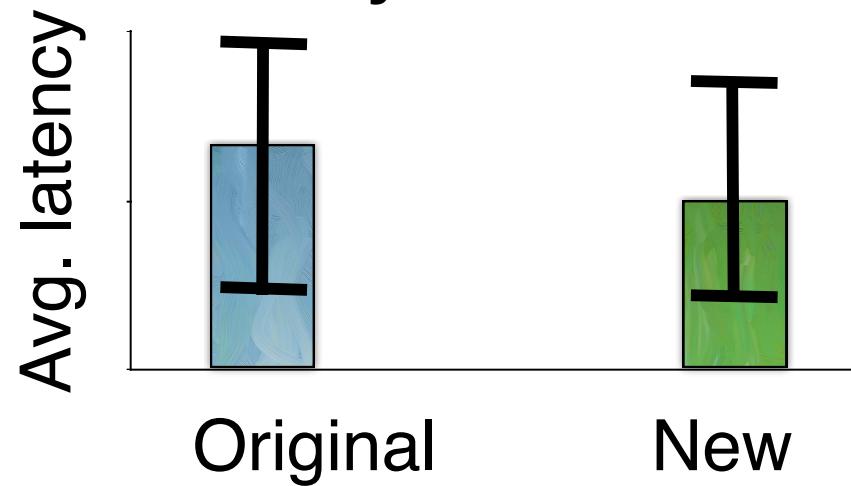


Image from Piled Higher & Deeper ([www.phdcomics.com](http://www.phdcomics.com))

# The easy life of a grad student

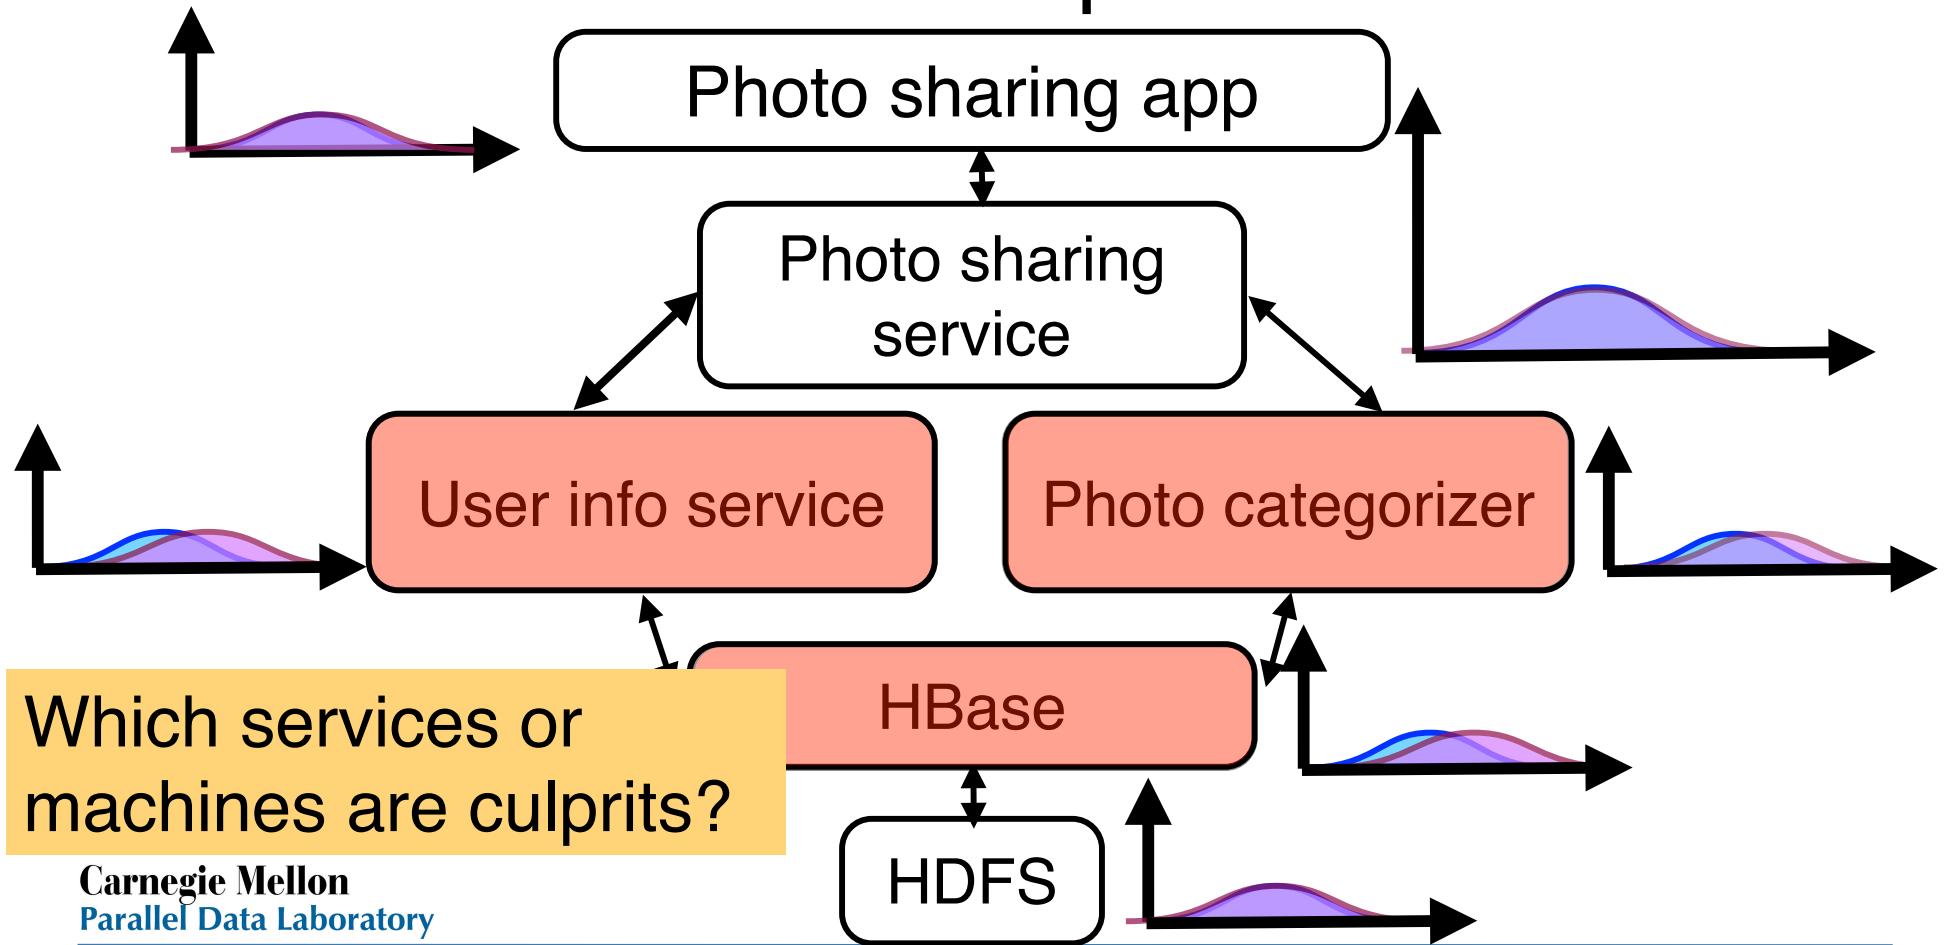
- Student's research idea improves avg. response time by 25%



- But, variance too high for confidence:
  - Can run more experiments
  - Can ignore variance ;)

# The hard life of a diagnosis tool

- Tool must find code responsible for an increase in mean response time



# The hard life of a diagnosis tool

---

- Can't ignore underlying variance
  - Will yield more false positives/negatives
  - FPs: Tool is misleading developer
  - FNs: Tool is not useful
- Can't 're-run' problem in production

Utility of diagnosis tools limited unless system has predictability properties

# How to increase predictability?

---

- Variance is treated as a key metric
  - Must be measured during dev/testing
- Devs must ID high-variance sources
  - Must *explicitly* decide on tradeoffs
- Must rigorously isolate irreducible variance sources from rest of system

# Grokking variance



(Developer)

## VarianceFinder

*Identifies similar behaviour w/high variance and localizes it*

## Variance nomenclature

*Helps devs reason about possible actions*

# The three I's of variance

- **Inadvertent** variance is unintentional

- E.g., spaghetti code or a bug

Reduce

- **Intrinsic** variance is fundamental

- E.g., Hard disks/non-networks

Isolate

- **Intentional** variance result of tradeoff

- E.g., low-latency schedulers

Isolate

# Open questions about variance

---

- How much needs to be reduced?
- If intrinsic variance is significant are hardware changes needed?
- If intentional variance is significant do we need to re-evaluate tradeoffs?

# Conclusion

---

- Predictability important for systems
- No answer but hard work during system development and testing
- Need tools/techniques to help developers increase predictability