

Vis  
2013

VAST • INFOVIS • SCIVIS  
BIOVIS • LDAV

Visualizing

request-flow comparison  
to aid performance diagnosis  
in distributed systems

Carnegie  
Mellon  
University

Raja Sambasivan,  
Ilari Shafer, Michelle Mazurek, Greg Ganger

# Distributed systems' traits & needs

- Used by society for many purposes



- Prone to difficult-to-solve problems
  - Due to scale & complexity
  - **Key need:** Sophisticated diagnosis tools

# Diagnosis technique/tool research

- Very active research area
  - E.g., Distalyzer, NetMedic, Spectroscope...
- Most only automatically localize problem
  - Developer must use results to ID root cause

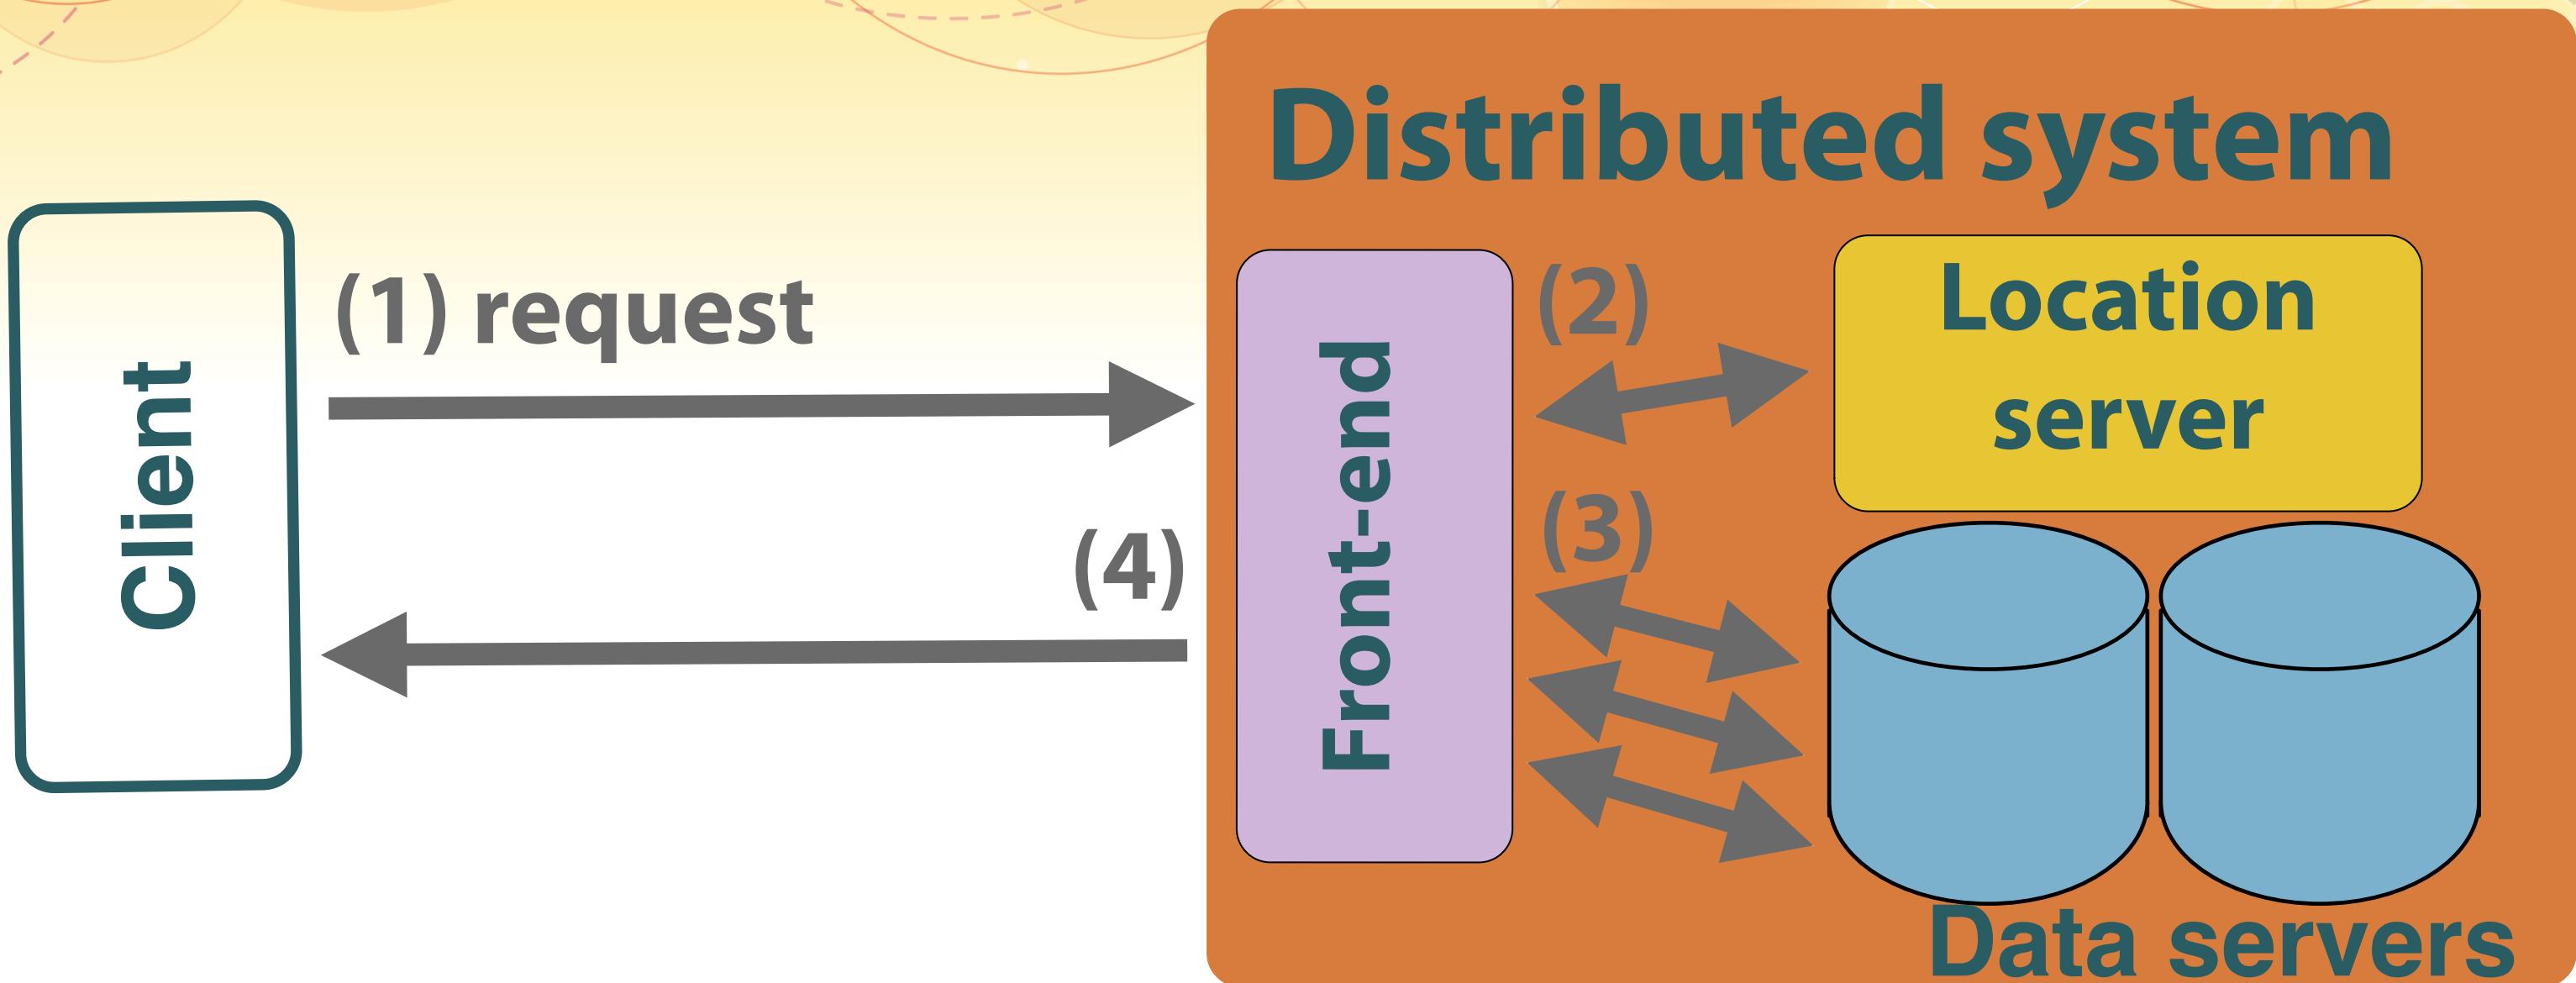
Tools must visualize results well to be effective

# Visualizing diagnosis tools' results

- Little guidance on what works best for **distributed system developers/problems**
  - There are a few exceptions
  - Limits utility of diagnosis tools

**This talk:** User studies to find good approaches for request-flow comparison

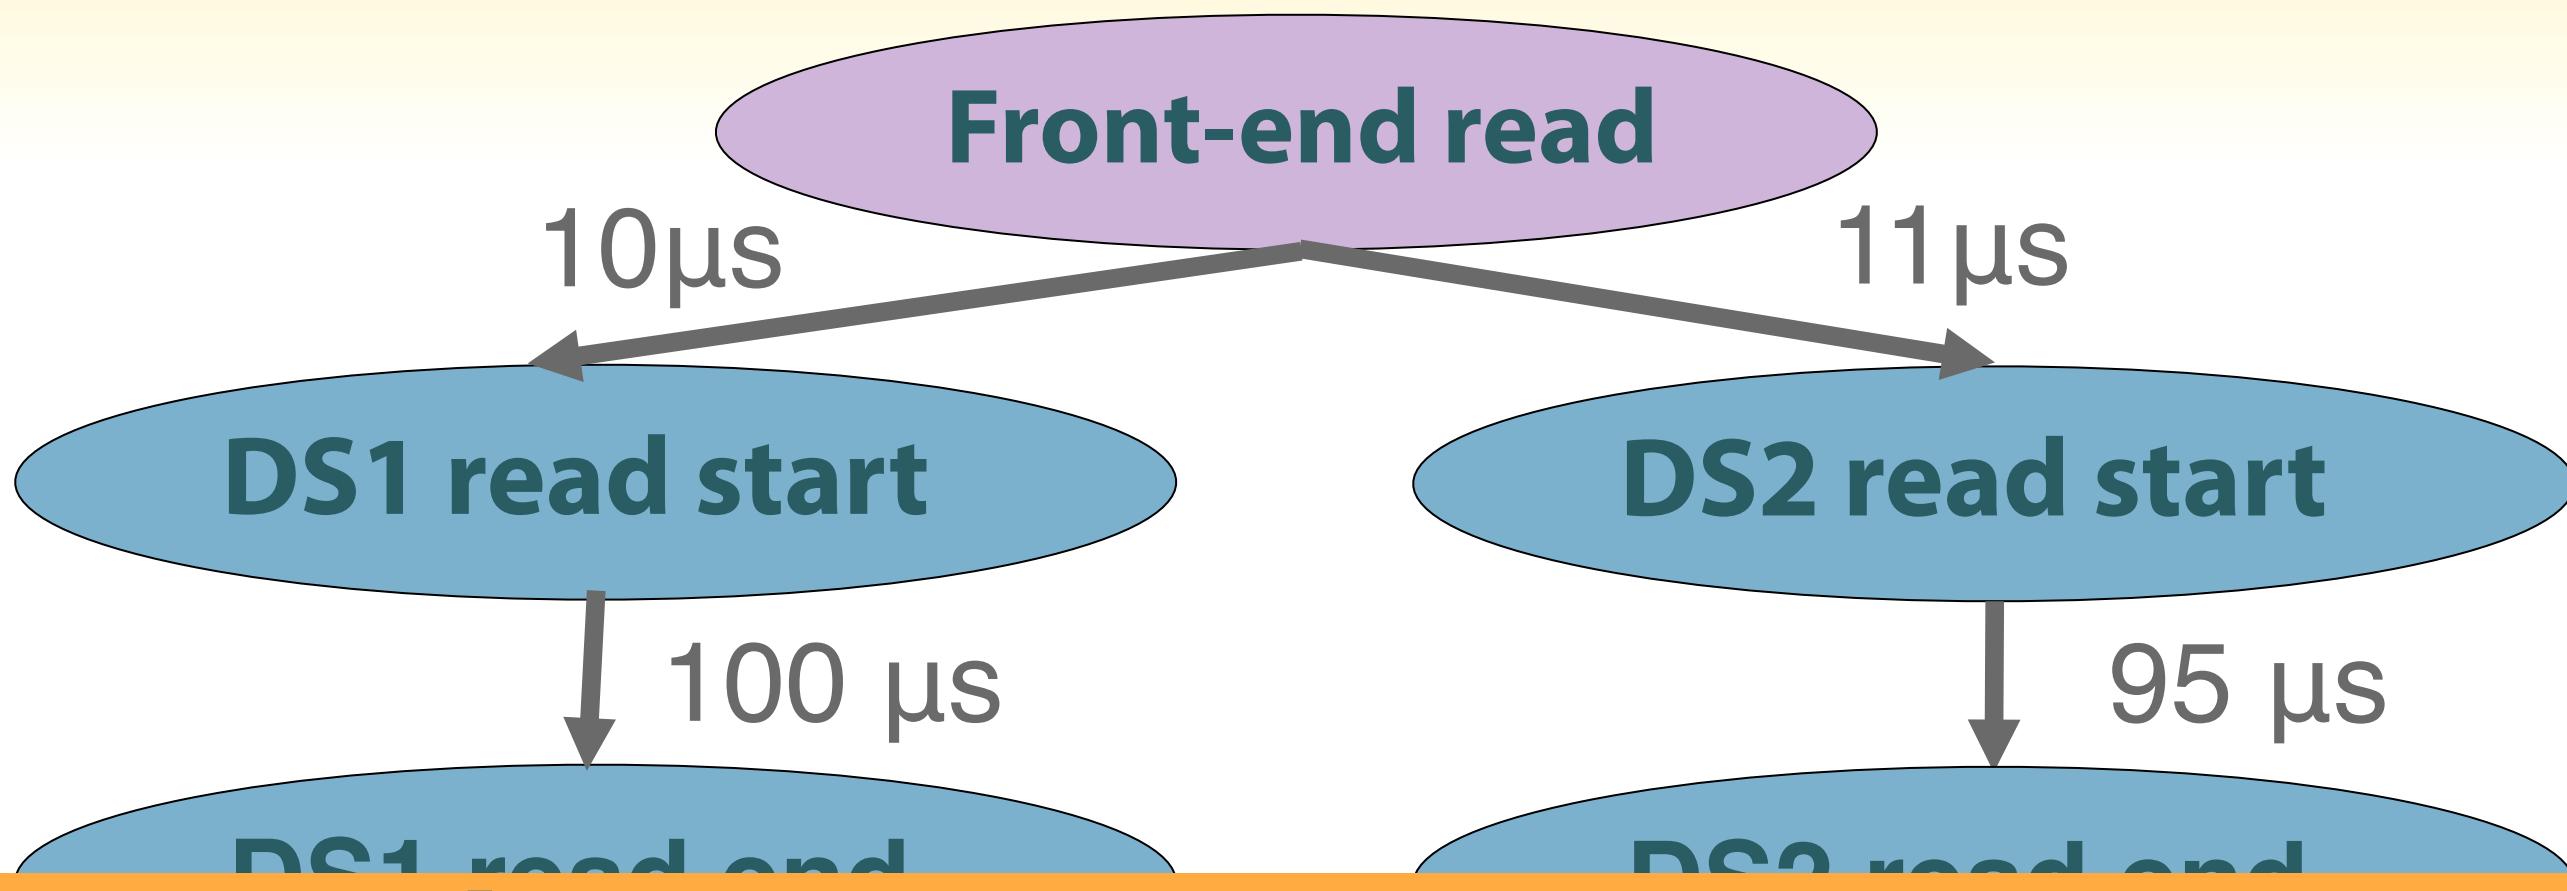
# Example distributed system



- Work done to satisfy request can be represented as a workflow (i.e., a DAG)

# An example request flow

Response time: 1,090µs



Nodes represent system events

Edges represent latency between events

# Spectroscope [NSDI'11]

- Localizes performance degradations
  - By ID'ing changed request flows
- Output:
  - Groups of before/after request flows
  - Some changes automatically ID'd

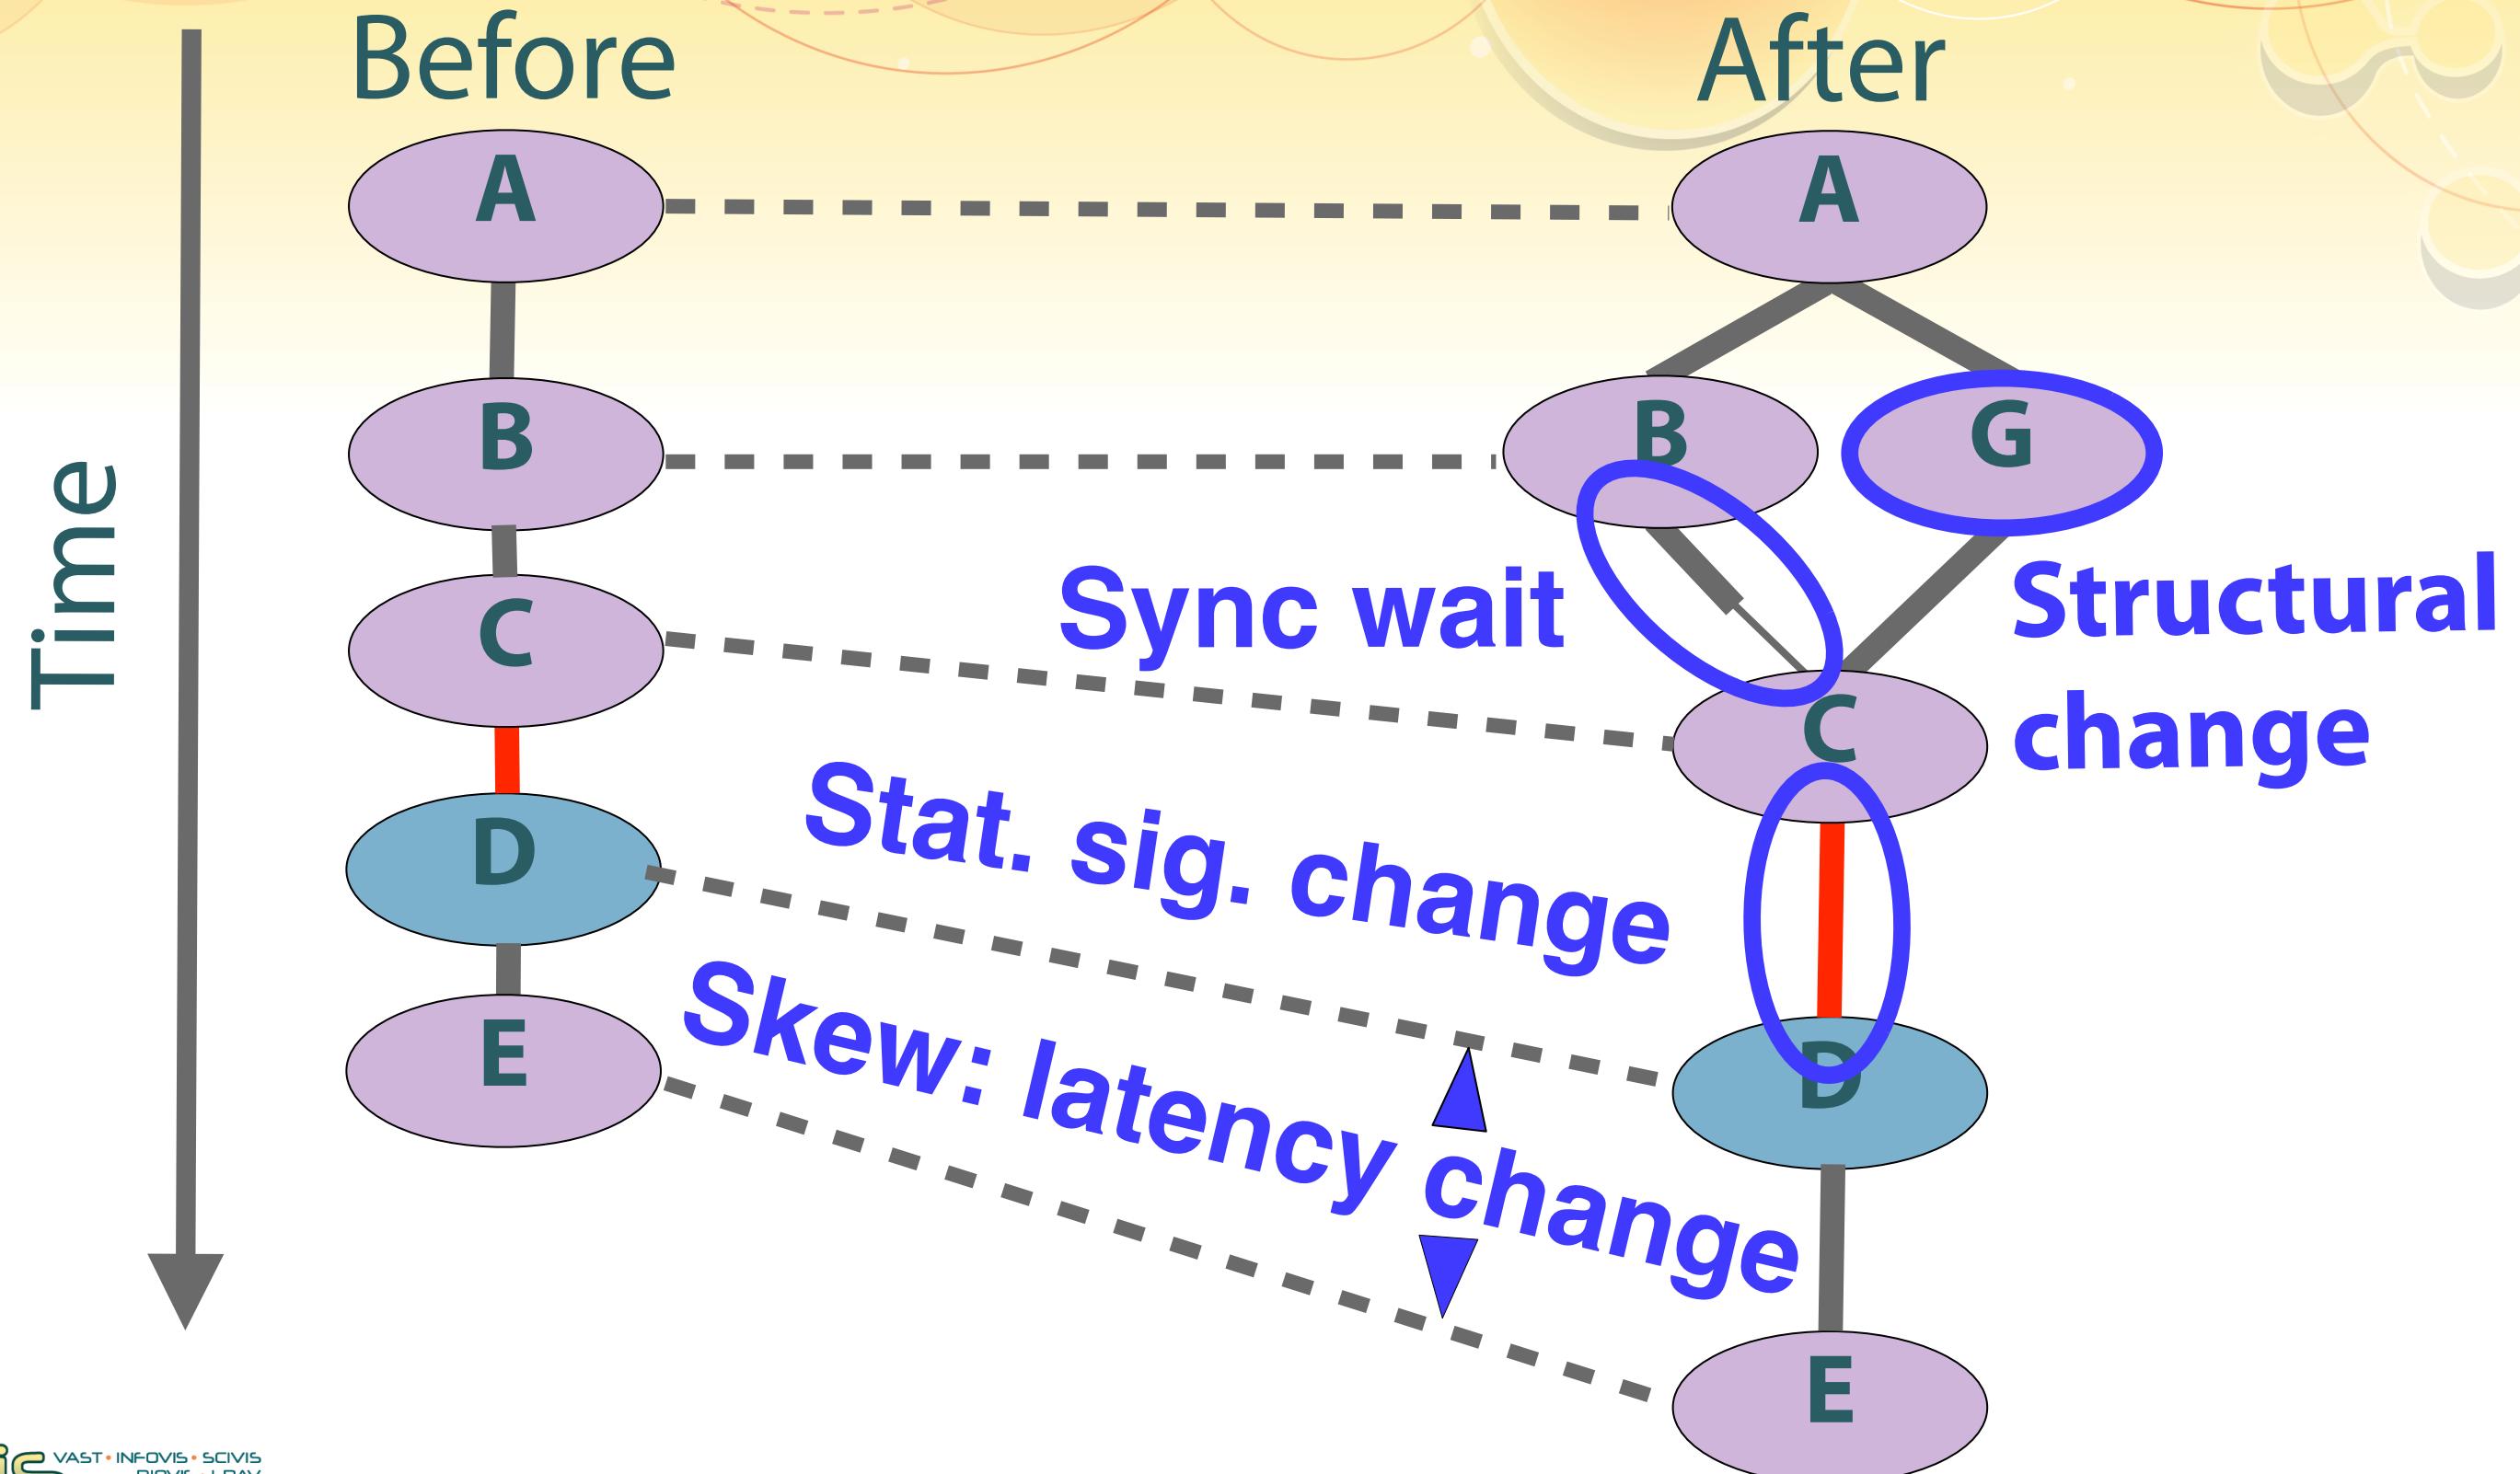
# Goals of visualization

- For a given before/after DAG in the output:
  - Help developers understand problem
  - Show automatically ID'd timing changes
  - Help identify important timing changes that weren't automatically identified
  - Help identify changed substructures

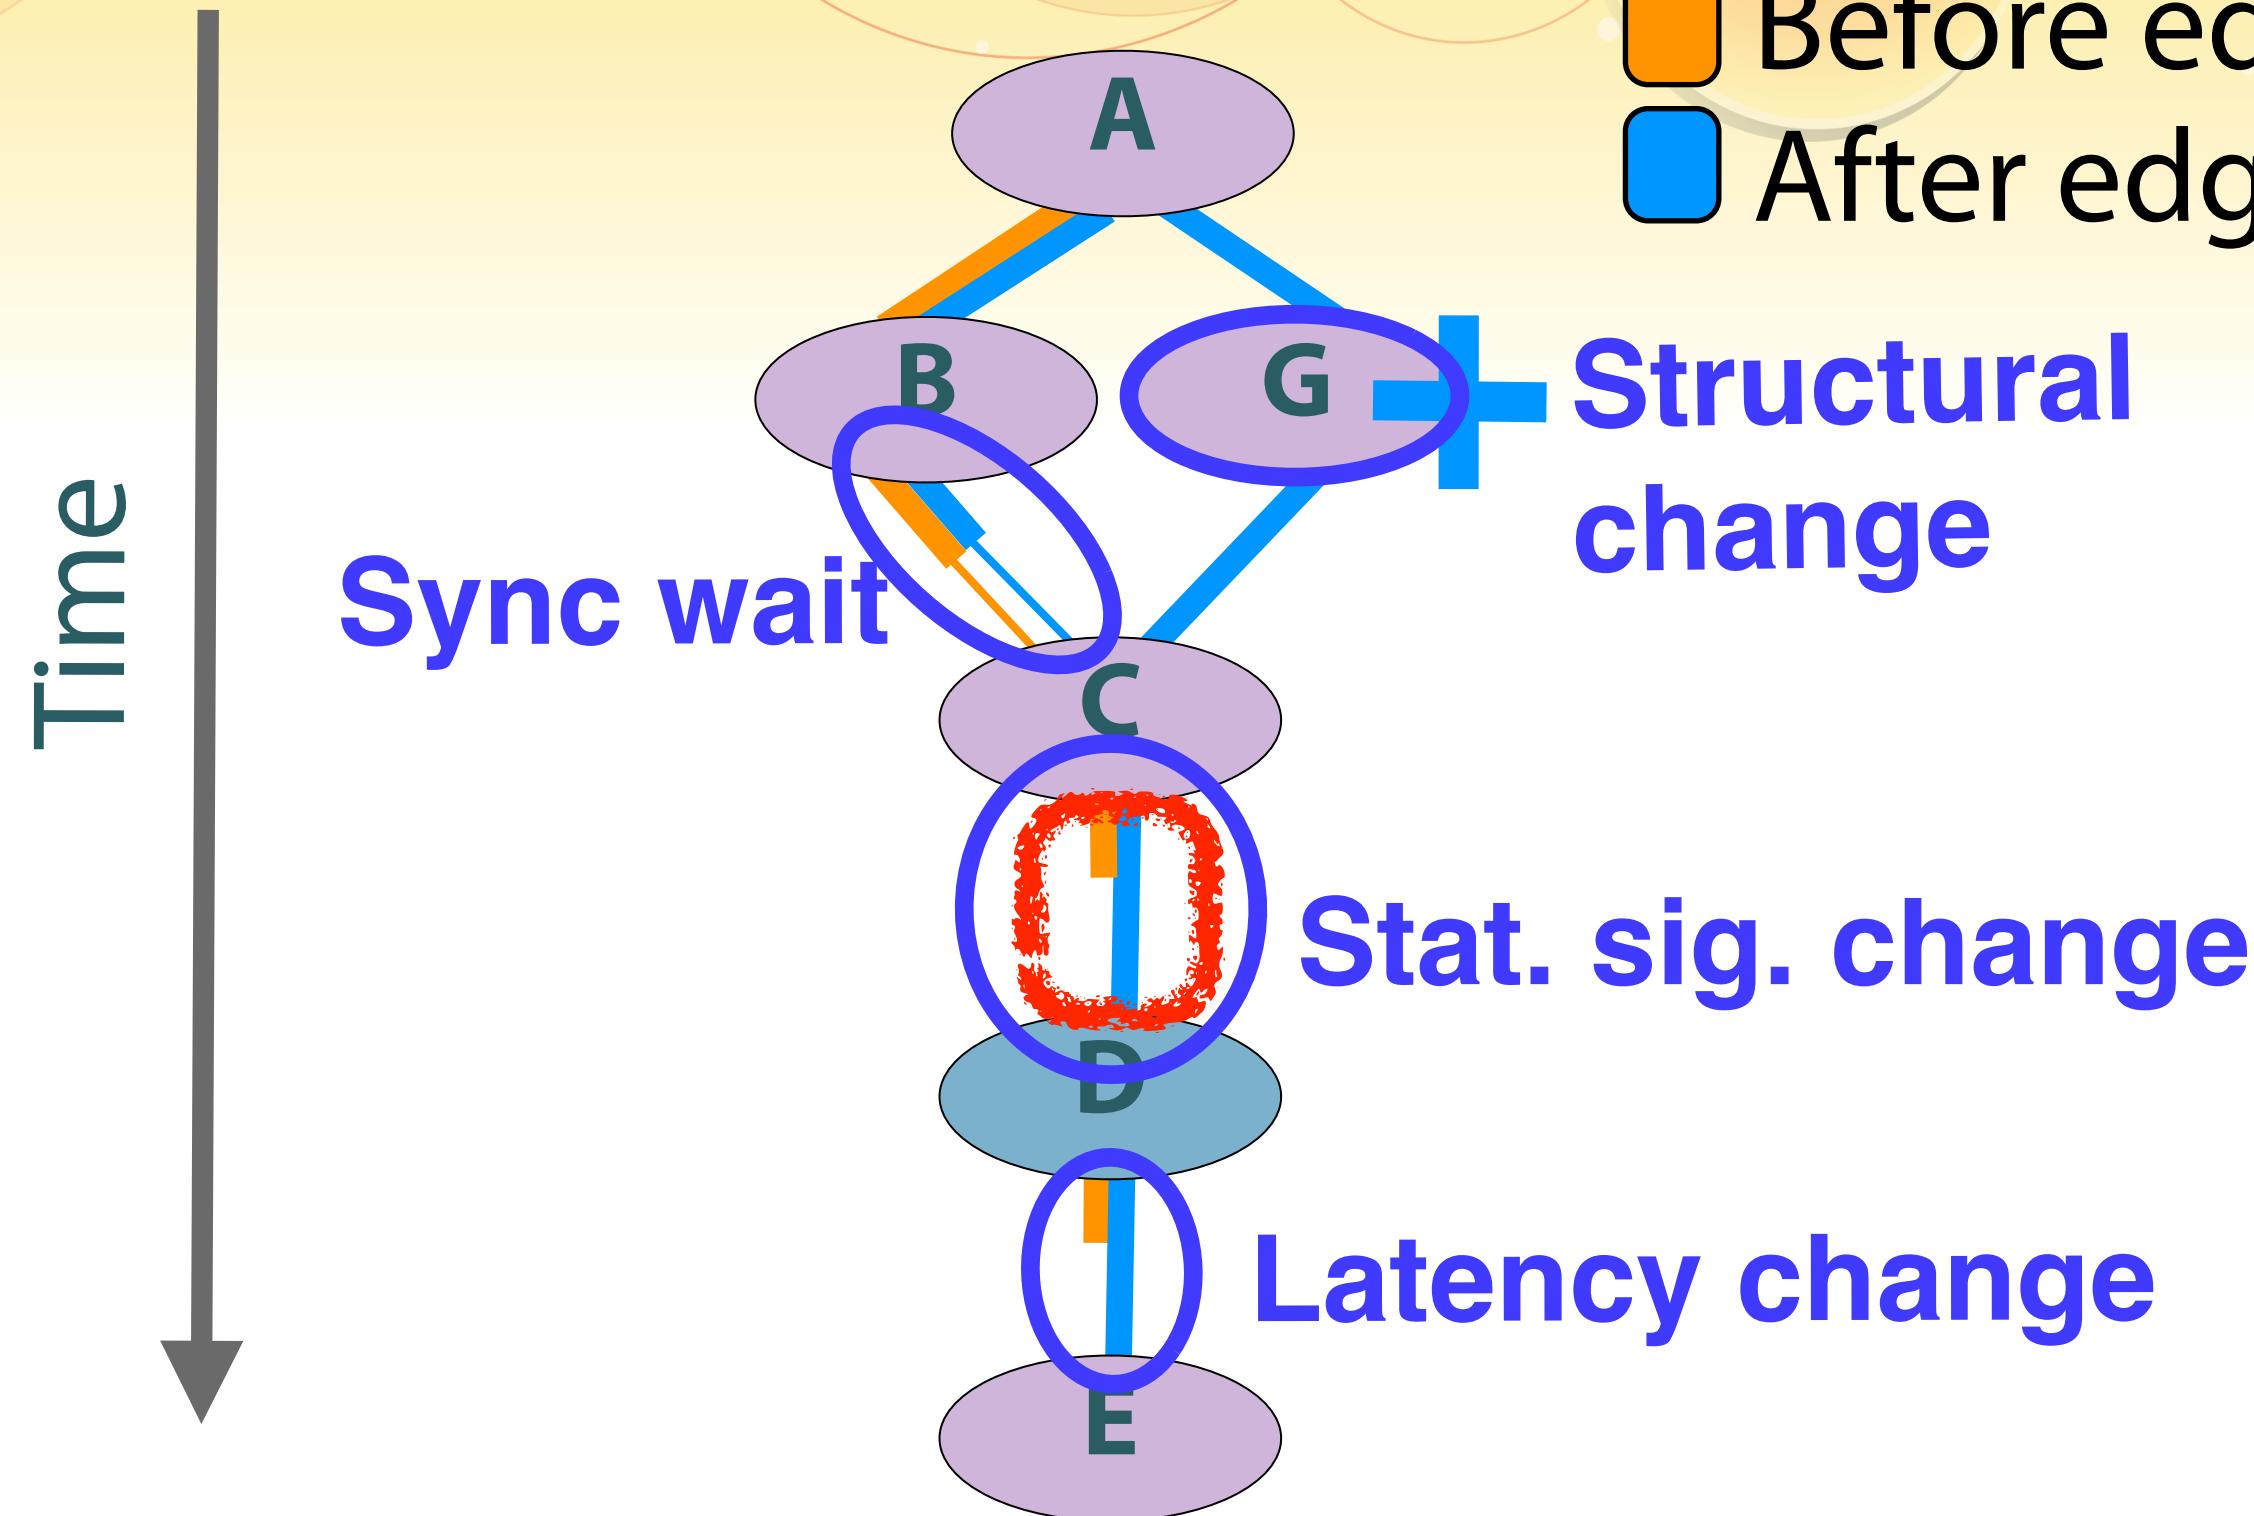
# Side-by-side, diff, animation

- Suggested by previous Spectroscope users who used a more basic presentation layer
- Represent orthogonal approaches
- Well-studied in different contexts

# Side-by-side implementation

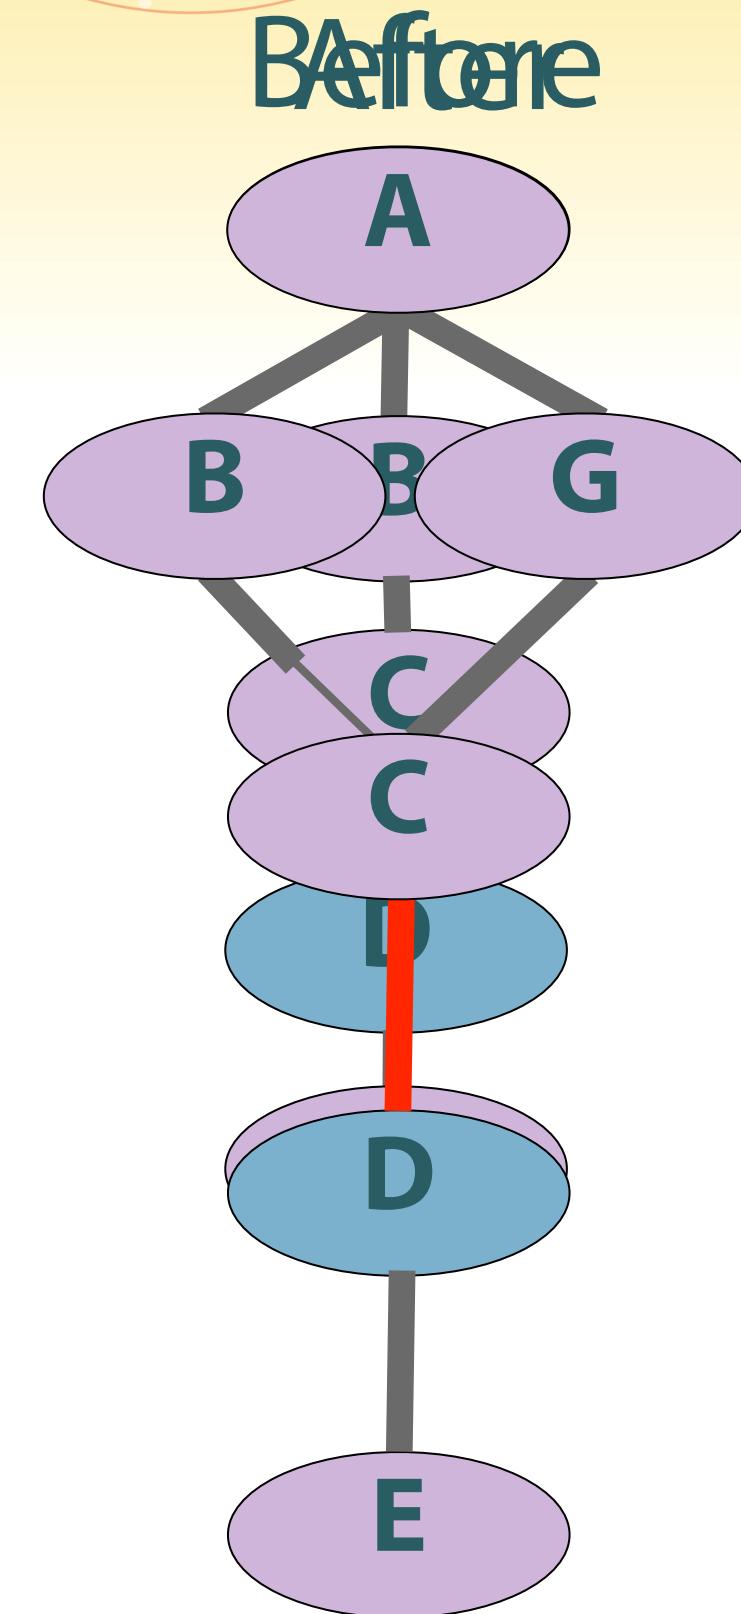


# Diff implementation



# Animation implementation

Time



before

after

▶ start

# User study overview & methodology

- Deployed a between subjects study
  - Composed of 26 participants
- Users diagnosed 5 Ursa Minor problems
  - Users asked to complete 5 tasks for each
  - Collected completion times & comments

# Overview of results

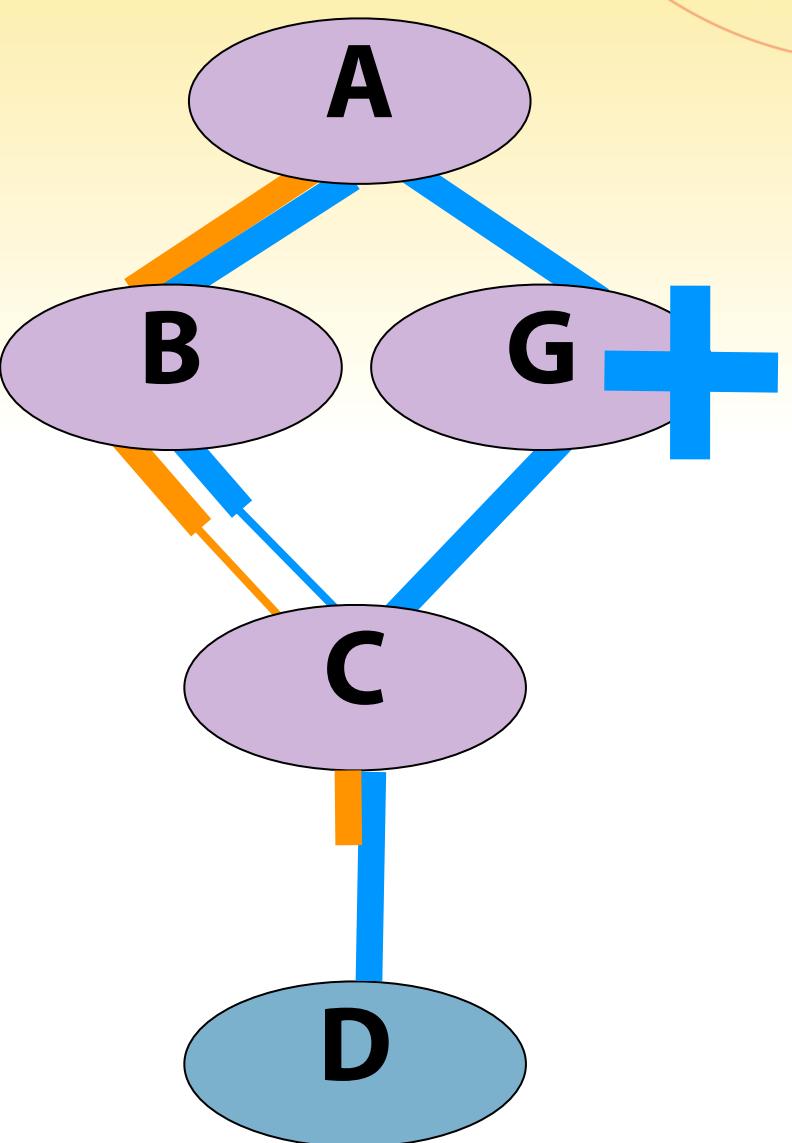
- No clear winner
  - If I had to choose...I'd be sad. (PD06)
- Animation approach slower for most tasks
- Different approaches best for diff. problems

# Best approach for problem types (I)

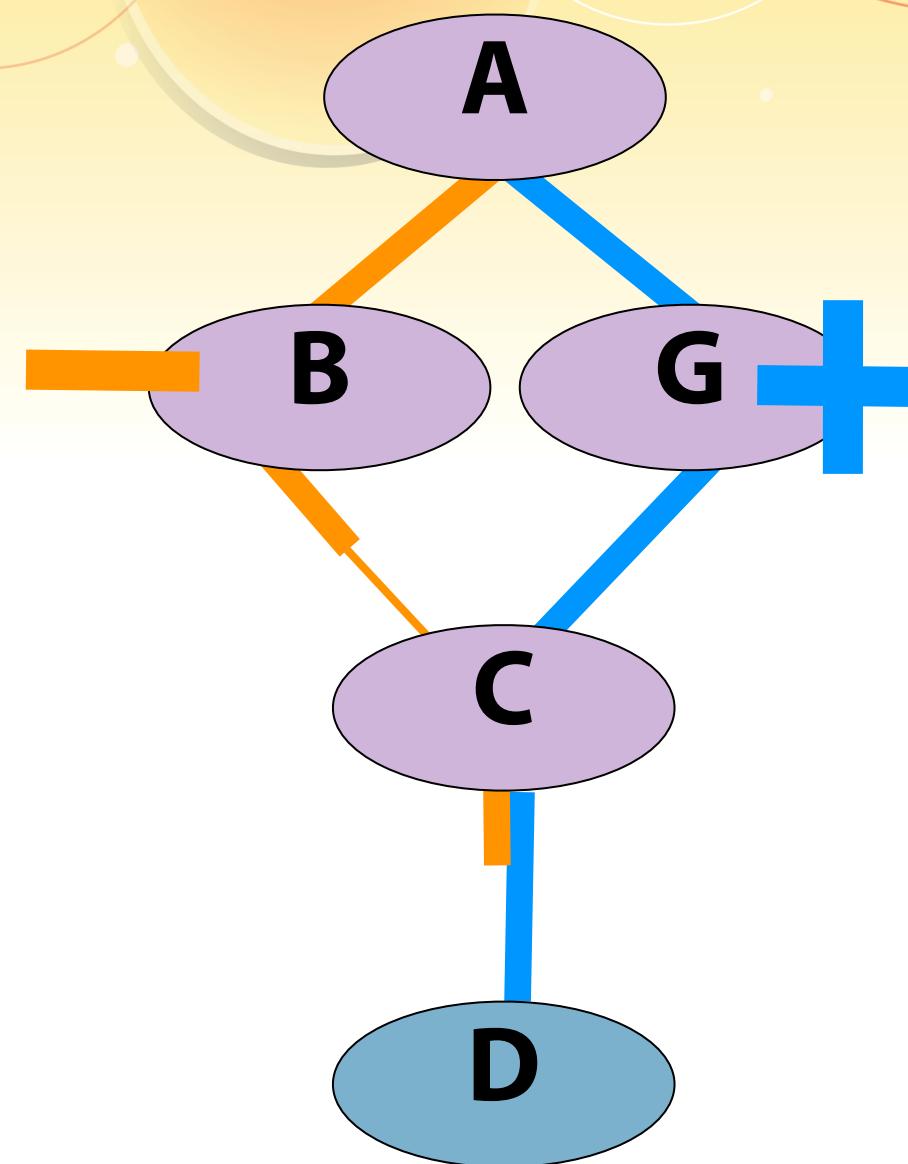
Approach	Edge latency	Concurrency	Intra-thread	Intra-thread mixed
Side Diff	✓			

# Diff: concurrency vs. intra-thread change

Extra concurrency



Intra-thread addition



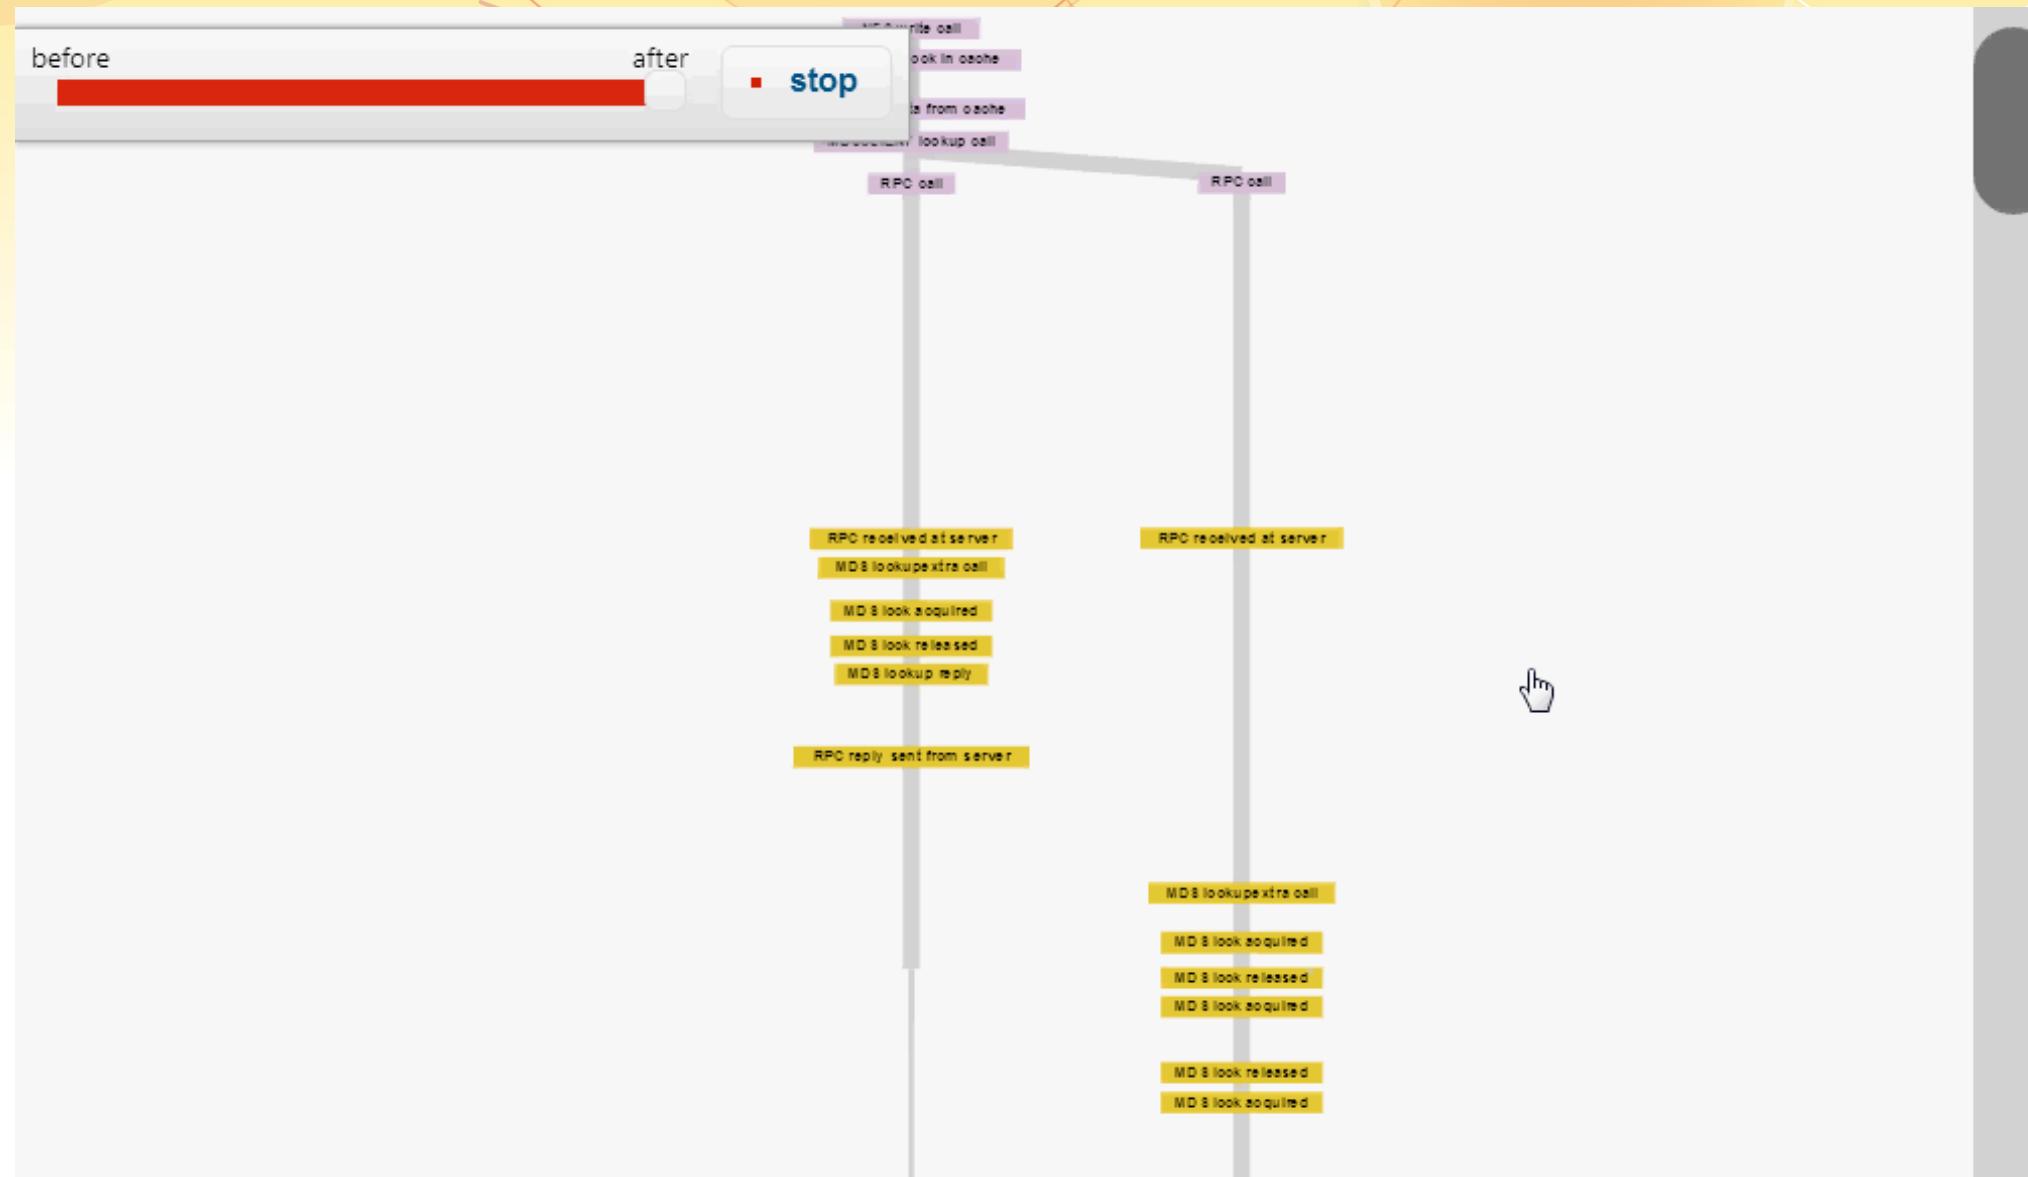
Similar representation

makes it easy to confuse one for the other

# Best approach for problem types (II)

Approach	Structural			
	Edge latency	Concurrency	Intra-thread	Intra-thread mixed
Side	✓			
Diff				
Anim		✓	✓	✗

# Anim: good for structural changes



Creates blinking effect that clearly demarcates differences and their impact

# Summary

- Conducted user study to find approaches suited to visualizing req.-flow comparison
- Found no single best approach
- See paper for more results & design lessons